

ERNIE Image — AI Image Generator with Perfect Text Rendering

13–16 minutes

Writing good prompts is the difference between a blurry approximation and a pixel-perfect AI-generated image. Since [ERNIE Image](#) open-sourced its 8B parameter Diffusion Transformer model in April 2026, creators worldwide have been discovering that this model rewards structured, deliberate prompting more than most of its competitors.

Unlike casual text-to-image tools where you type a sentence and hope for the best, ERNIE Image responds exceptionally well to methodical prompt construction — and its built-in Prompt Enhancer means you can start simple and iterate toward complexity.

This guide covers everything you need to know about writing effective prompts for ERNIE Image, from fundamental syntax patterns to advanced techniques for structured layouts, multilingual text rendering, and style control.





Before diving into prompt patterns, it helps to understand what happens under the hood. ERNIE Image uses a two-stage prompt pipeline:

Stage 1 — Prompt Enhancer (optional): A lightweight 3B-parameter language model (based on Ministral-3B) automatically expands your input into a more detailed, structured description. If

you write "a cat on a windowsill," the Enhancer might expand it to "a fluffy orange tabby cat sitting gracefully on a wooden windowsill, warm afternoon sunlight streaming through sheer curtains, soft bokeh background of a cozy living room."

Stage 2 — Diffusion Transformer: The enhanced (or raw) prompt is processed by the 8B DiT model to generate the image through 50 denoising steps (or 8 steps with the Turbo variant).

This two-stage architecture is why ERNIE Image works well with both short and long prompts — but understanding how to leverage each stage gives you far more control over the output.

A well-structured ERNIE Image prompt typically contains four components:

`[Subject] + [Description/Details] + [Style/Medium] + [Technical Parameters]`

Not every prompt needs all four, but addressing each deliberately produces more predictable results.

[Component 1: Subject](#)

Be specific. Instead of "a dog," write "a golden retriever puppy."
Instead of "a building," write "a Victorian brownstone townhouse."

Weak: a landscape **Strong:** a terraced rice field carved into misty mountains in Bali

[Component 2: Description and Details](#)

Add sensory and spatial details. ERNIE Image has strong instruction-following capabilities (it scored 0.8856 on the GENEval benchmark, leading all open-weight models), so specific positioning, lighting, and atmosphere descriptions translate well.

Weak: a woman in a garden **Strong:** a woman in her thirties wearing a linen sundress, standing among blooming lavender bushes, golden hour sunlight casting long shadows

[Component 3: Style and Medium](#)

ERNIE Image supports an impressive range of styles. Explicitly stating the desired style prevents ambiguity:

- **Photography:** "shot on Hasselblad, 85mm lens, shallow depth of field, natural lighting"
- **Illustration:** "vector illustration, flat design, limited color palette, minimalist"
- **Cinematic:** "cinematic still, anamorphic lens flare, teal and orange color grading"
- **Watercolor:** "watercolor painting on textured paper, wet-on-wet technique, soft edges"
- **3D Render:** "octane render, volumetric lighting, subsurface

scattering, PBR materials"



[Component 4: Technical Parameters](#)

When using the API or local deployment, you can set resolution, inference steps, and guidance scale directly. When prompting

through a web interface, you can still influence output quality by including terms like "high resolution, sharp focus, 4K, detailed textures."

The Prompt Enhancer is ERNIE Image's secret weapon, but it's not always the right choice.

[When to enable Prompt Enhancer \(use_pe=True\)](#)↻

- You're starting with a short, simple concept and want the model to fill in details
- You're new to prompt engineering and want better results with less effort
- You're generating structured content like posters, infographics, or comic panels
- You want bilingual text rendering (the Enhancer handles mixed-language text placement well)

[When to disable Prompt Enhancer \(use_pe=False\)](#)↻

- You've already crafted a detailed, precise prompt and don't want it altered
- You need exact control over text rendering in the image
- You're iterating on a specific composition and changes to the prompt would disrupt your workflow

- The GENEval benchmark actually shows slightly higher scores without the Enhancer (0.8856 vs 0.8728), meaning raw precision is sometimes better

Based on extensive testing and community feedback, here are proven prompt patterns for common use cases.

[Pattern 1: Photorealistic Portraits](#)

A portrait of [subject], [clothing/accessories],
[expression], [lighting],
[background], shot on [camera/lens], [post-processing
style]

Example: "A portrait of an elderly Japanese man with deep wrinkles and kind eyes, wearing a dark blue kimono, gentle smile, soft window light from the left, blurred tatami room interior, shot on Sony A7III 85mm f/1.4, warm color tones, film grain"

[Pattern 2: Product and Marketing Images](#)

A [product type] on [surface/background], [lighting
setup], [color scheme],
[additional props], commercial photography style,
[mood]

Example: "A sleek ceramic pour-over coffee maker on a raw concrete surface, warm directional spotlight from upper right, earth

tones with brass accents, scattered coffee beans and a linen napkin, commercial photography style, clean and modern"

[Pattern 3: Illustrated Infographics](#)

For structured layouts, be explicit about the grid and text placement. As covered in our [creating posters and comics guide](#), ERNIE Image excels at multi-cell layouts:

```
An infographic poster about [topic], divided into [N]
sections arranged in
[layout description], each section contains [content
description], [style],
[color scheme], title text "[Title]" at the top
```

Example: "An infographic poster about the solar system, divided into 8 sections arranged in 2 rows of 4, each section contains a planet illustration with its name and key facts, flat vector illustration style, dark navy background with white and gold text, title text 'Our Solar System' at the top"

[Pattern 4: Comic and Manga Panels](#)

ERNIE Image's text rendering capability (scoring 0.9733 on LongTextBench) makes it uniquely suited for comic generation:

```
A [N]-panel comic page, [style/manga style], panel 1:
[description],
```



```
panel 2: [description], ... speech bubbles containing  
"[dialogue]",  
[overall mood/coloring]
```

[Multilingual Text Rendering Tips](#)

One of ERNIE Image's biggest advantages over competitors like FLUX — which completely collapses on Chinese text (0.2183 vs ERNIE's 0.9661) — is its bilingual text rendering. Here's how to get the best results:

1. **Wrap text in quotes:** To render specific text in an image, wrap the target text in double quotation marks within your prompt
2. **Specify the language:** If mixing languages, explicitly state which text is in which language
3. **Use the Prompt Enhancer for mixed scripts:** The Enhancer handles mixed Chinese-English layouts particularly well
4. **Keep text short per element:** Each text element should be under 10-15 characters for best rendering accuracy

Example prompt with bilingual text: "A modern coffee shop menu board, written in both English and Chinese, left column shows 'Espresso / 浓缩咖啡', right column shows 'Latte / 拿铁', chalkboard style, warm ambient lighting"

[Using Art History References](#)

Naming specific art movements or artists helps anchor the style:

- "Art Nouveau style, inspired by Alphonse Mucha, flowing organic lines, decorative borders"
- "Bauhaus design, geometric shapes, primary colors, sans-serif typography"
- "Ukiyo-e woodblock print style, flat colors, bold outlines, wave patterns"

[Using Mood and Atmosphere Keywords](#)

ERNIE Image is highly responsive to emotional and atmospheric cues:

- "Melancholic, overcast day, desaturated colors, lonely figure"
- "Vibrant and energetic, neon lights, rain-soaked streets, cyberpunk"
- "Serene and meditative, foggy morning, muted earth tones, zen garden"

[Combining Style References](#)

You can layer multiple style references for unique results:

"Photography influenced by Gregory Crewdson's cinematic staging, combined with Edward Hopper's use of light and shadow, a diner scene at 2 AM"

Mistake 1: Overloading with Contradictory Instructions

Problem: "A photorealistic watercolor painting in 3D render style with anime characters" **Fix:** Pick one primary style and one secondary influence: "A watercolor painting with subtle 3D volumetric lighting effects"

Mistake 2: Vague Spatial Descriptions

Problem: "Objects arranged nicely" **Fix:** "Five objects arranged in a horizontal row, evenly spaced, from left to right: a red apple, a blue book, a green plant, a yellow mug, and a white candle"

Mistake 3: Ignoring the Prompt Enhancer's Impact

Problem: Spending 30 minutes crafting a perfect prompt, then leaving the Enhancer on and having it rewrite your work **Fix:** Toggle the Enhancer based on your needs. If you've already written a 50+ word prompt with specific details, turn it off.

Mistake 4: Not Specifying Aspect Ratio

Problem: Always generating square images regardless of content needs **Fix:** ERNIE Image supports multiple resolutions — 1024x1024 (square), 1264x848 (landscape), 848x1264 (portrait), and more. Match the aspect ratio to your content. Panoramic landscapes need wide formats; portraits need tall formats.

The best results come from systematic iteration. Here's a recommended workflow:

1. **Start with the Turbo model** (8 inference steps, guidance_scale=1.0) for rapid prototyping
2. **Write a basic prompt** and enable the Prompt Enhancer to see how the model interprets your concept
3. **Refine the prompt** based on the Turbo output — add details the model missed, adjust style references
4. **Disable the Enhancer** once your prompt is detailed enough
5. **Switch to the Standard model** (50 inference steps, guidance_scale=4.0) for the final render
6. **Fine-tune resolution** to match your intended use case

For more on setting up both models locally, see our [local installation guide](#).

[Technique 1: Negative Prompting via Exclusion](#) ↻

While ERNIE Image doesn't have a formal negative prompt parameter in the same way Stable Diffusion does, you can achieve similar effects by explicitly describing what you want to avoid within the positive prompt:

"A clean minimalist living room, no clutter, no visible cables, no people, strictly architectural photography"

[Technique 2: Reference-Based Prompting](#)

Reference specific works, films, or visual traditions to anchor the aesthetic:

"In the visual style of Studio Ghibli's Spirited Away, a bathhouse entrance at twilight, lanterns glowing, spirits gathering"

[Technique 3: Step-Back Prompting for Complex Scenes](#)

For complex multi-element scenes, describe the overall composition first, then layer in details:

"A vast underground cathedral made of salt crystals — massive stalactites hanging from the ceiling refracting prismatic light, a still underground lake reflecting the crystal formations, tiny human figure at the bottom for scale, fantasy art style, volumetric god rays"

[Technique 4: Leveraging the Prompt Library](#)

If you're looking for inspiration, check out our [curated ERNIE Image prompt collection](#) with tested prompts organized by category — from photorealistic scenes to stylized illustrations to structured layouts.

When using ERNIE Image through APIs (WaveSpeed AI, fal.ai, Replicate, or self-hosted), consider these additional tips:

1. **Consistent encoding:** Ensure your prompt text is UTF-8 encoded,

especially for Chinese and other non-Latin scripts

2. **Prompt length:** While ERNIE Image handles long prompts well, the sweet spot is 50-150 words for optimal results
3. **Temperature and guidance:** Use `guidance_scale=4.0` for the Standard model and `guidance_scale=1.0` for Turbo — these aren't arbitrary but optimized values from the ERNIE team
4. **Batch prompting:** When generating series of images, keep the style description consistent across prompts and only vary the subject content

According to the [official ERNIE-Image GitHub repository](#), the model was extensively beta-tested across 30+ enterprise partners and refined based on professional creator feedback, which means these parameter recommendations are well-validated.

Use Case	Template
Portrait	"[Subject], [clothing], [expression], [lighting], [background], shot on [camera], [style]"
Landscape	"[Location/scene], [time of day], [weather/atmosphere], [camera angle], [style]"
Product	"[Product] on [surface], [lighting], [props], commercial photography, [color scheme]"
Poster	"A poster for [purpose], [layout], [text: content],

Use Case	Template
	[style], [color scheme]"
Comic	"[N]-panel comic, [style], panel descriptions, speech bubbles, [coloring]"
Infographic	"Infographic about [topic], [layout/grid], [data sections], [style], [colors]"
Concept Art	"[Subject], concept art, [environment], [mood], [style reference], detailed, [lighting]"

ERNIE Image rewards deliberate prompt engineering more than most text-to-image models. Its dual-stage architecture — the Prompt Enhancer for quick exploration and the raw DiT for precise control — gives you two distinct workflows in one model. The key takeaways:

- **Start simple** with the Enhancer enabled, then iterate toward precision
- **Be explicit** about style, layout, and spatial relationships
- **Leverage bilingual text** — it's ERNIE Image's most distinctive capability
- **Use Turbo for drafts**, Standard for finals
- **Match your aspect ratio** to your content type

With practice, you'll develop an intuition for how ERNIE Image interprets different prompt structures. The model is still young in the open-source community, and prompt techniques will continue to evolve as more creators experiment with it. For a constantly updated collection of tested prompts, bookmark our [ERNIE Image prompt library](#).

Happy generating.